

# OpenGuardrails: An Open-Source Context-Aware AI Guardrails Platform

Thomas Wang<sup>1</sup>      Haowen Li<sup>2</sup>  
<sup>1</sup> OpenGuardrails.com    <sup>2</sup> The Hong Kong Polytechnic University

October 20, 2025

<https://openguardrails.com>  
<https://github.com/openguardrails/openguardrails>  
<https://huggingface.co/openguardrails>

## Abstract

As large language models (LLMs) become increasingly integrated into real-world applications, safeguarding them from unsafe, malicious, or privacy-violating content is of critical importance. We present **OpenGuardrails**, the first open-source project that provides both a *context-aware safety & manipulation-detection model* and a *deployable platform* for comprehensive AI guardrails. OpenGuardrails protects against content-safety risks, model-manipulation attacks (e.g., prompt injection, jailbreaking, code-interpreter abuse and generation/execution of malicious code), and data leakage. Content-safety and model-manipulation detection are implemented by a unified large model; data-leakage identification and redaction are performed by a separate lightweight NER pipeline (e.g., presidio-style models or regex-based detectors). The system can be deployed as a security gateway or API-based service with enterprise-grade, fully private deployment options. OpenGuardrails achieve state-of-the-art (SOTA) performance on safety benchmarks, excelling in both prompt and response classification across English, Chinese, and multilingual tasks. All models are released under the Apache 2.0 license for public use.

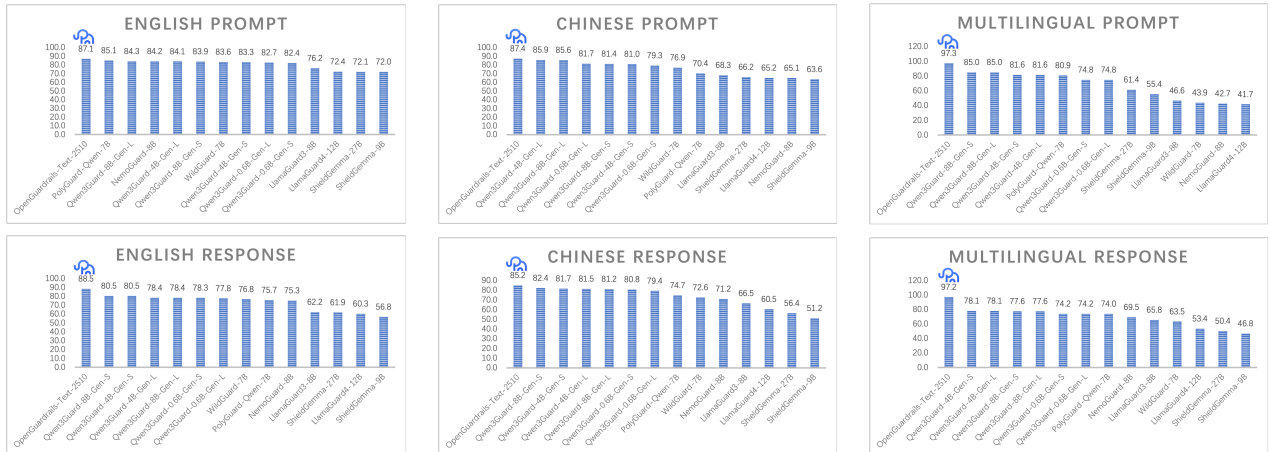


Figure 1: Average F1 scores of OpenGuardrails vs. existing guard models across safety classification benchmarks for Prompts and Responses in English, Chinese, and Multilingual datasets.

# 1 Introduction

Recent advances in foundation models such as GPT-5, Claude 4, Gemini 2.5, and Qwen3 have enabled powerful capabilities in language understanding and generation. However, these capabilities introduce several classes of risks: (1) *content-safety* risks where the model produces harmful, hateful, illegal, or sexually explicit content; (2) *model-manipulation attacks* such as prompt injection, jailbreaks, and code-interpreter abuse that attempt to trick a model into generating or executing malicious code; and (3) *data leakage* where sensitive personal or organizational information is exposed.

**OpenGuardrails** is designed to address this full risk spectrum in a unified, deployable system. Unlike prior efforts that open-source either models or rule-based tools in isolation, OpenGuardrails provides:

- **A unified large model** for both content-safety and model-manipulation detection.
- **A separate lightweight NER/data-redaction pipeline** for identifying and masking sensitive information.
- **A production-ready platform** that exposes APIs and deployment scripts for enterprise integration.

**The main contributions of OpenGuardrails are as follows:**

1. **Configurable Policy Adaptation.** We propose a practical solution to the *policy inconsistency* problem commonly observed across safety benchmarks and guard models. Inspired by the findings in the Qwen3Guard report, we recognize that inconsistent annotation criteria—such as those in ToxicChat and OpenAIModeration—lead to unreliable evaluations. Instead of introducing a “Controversial” label (which increases human review costs in production), OpenGuardrails allows users to *configure unsafe categories* and *set sensitivity thresholds* based on their application domain. The guard model outputs either “safe” or “unsafe” along with a confidence score derived from the probability of the first token. Administrators can tune the system’s sensitivity (high, medium, or low) so that higher sensitivity trusts lower-confidence unsafe signals, enabling automated, flexible control over safety tolerance in enterprise environments.
2. **Unified LLM-based Guard Architecture.** OpenGuardrails demonstrates that a single LLM can perform production-grade content-safety and model-manipulation detection. Compared with prior systems such as *LlamaFirewall*, which relies on PromptGuard 2 (a fine-tuned BERT-style classifier), our unified LLM approach provides superior semantic understanding of complex attacks and simplifies deployment pipelines.
3. **Fully Open-source and Deployable System.** OpenGuardrails is the first to release both a large guard model and a production-ready platform, enabling users to deploy privately within their infrastructure or extend the system for custom development.
4. **Scalable and Efficient Model Design.** The guard model is fine-tuned from a 14B dense base model and quantized via GPTQ to 3.3B parameters, achieving high throughput and concurrency suitable for real-time production use. Prior works rarely scale beyond 8B models while maintaining comparable latency efficiency.
5. **Multilingual Safety Coverage.** OpenGuardrails supports 119 languages and dialects, ensuring robust performance for global and cross-lingual applications.
6. **State-of-the-art Performance.** The system achieves SOTA results on multilingual safety benchmarks, excelling in both prompt and response classification tasks.

7. **Open Data Contribution.** We release *OpenGuardrailsMixZh\_97k*, a collection of five translated Chinese safety datasets—ToxicChat, WildGuardMix, PolyGuard, XSTest, and BeaverTails—under the Apache 2.0 license on HuggingFace, totaling 97k samples to promote multilingual safety research.

## 2 Contributions Summary

We summarize the main contributions of **OpenGuardrails** as follows:

1. **Configurable Safety Policy Mechanism.** OpenGuardrails introduces a practical and enterprise-oriented solution to the long-standing *policy inconsistency* problem observed in existing safety benchmarks and guard models. Instead of relying on human-validated “Controversial” labels as proposed in Qwen3Guard, our approach allows users to *configure unsafe categories* and *adjust sensitivity thresholds* dynamically. The guard model outputs probabilistic confidence scores, enabling administrators to tune detection sensitivity (high, medium, low) according to their risk tolerance and application context. This mechanism provides a cost-effective and adaptive framework for enterprise-grade safety governance.
2. **Unified LLM-based Guard Architecture.** We demonstrate that a single large language model can effectively perform both content-safety detection and model-manipulation defense (e.g., prompt injection, jailbreaking, and code-interpreter abuse). Compared with hybrid architectures like *LlamaFirewall*, which depend on smaller BERT-style detectors, OpenGuardrails achieves superior semantic understanding, robustness, and ease of deployment.
3. **First Fully Open-source Guardrail System.** OpenGuardrails is the first project to open-source both a large-scale safety LLM and a production-ready guardrail platform. It provides API interfaces, deployment scripts, and modular components for seamless private or on-premise deployment, empowering enterprises and researchers to build upon a transparent and extensible safety infrastructure.
4. **Scalable and Efficient Model Design.** The OpenGuardrails-Text-2510 model is fine-tuned from a 14B dense base model and quantized via GPTQ to 3.3B parameters. This configuration achieves low latency ( $P95 = 274.6$  ms) and high throughput suitable for real-time applications, while maintaining state-of-the-art accuracy across safety benchmarks. Prior open-guard systems rarely scale beyond 8B while preserving production-level efficiency.
5. **Multilingual and Cross-domain Coverage.** OpenGuardrails supports 119 languages and dialects, providing reliable safety coverage across global and multilingual contexts. It consistently achieves top-tier performance on English, Chinese, and multilingual benchmarks in both prompt-level and response-level classification tasks.
6. **Open Safety Data Contribution.** We release *OpenGuardrailsMixZh\_97k*, a new Chinese safety dataset collection which includes translated and aligned versions of ToxicChat, WildGuardMix, PolyGuard, XSTest, and BeaverTails. This dataset (97k samples) is publicly available on HuggingFace under the Apache 2.0 license to facilitate further research in multilingual safety evaluation and model alignment.

## 3 Related Work

Large language model (LLM) safety research has rapidly evolved in recent years, focusing on mitigating harmful content, preventing manipulation attacks, and reducing private data leakage. In this section, we

compare **OpenGuardrails** against major prior guard frameworks, including *Qwen3Guard*, *LlamaFirewall*, *PromptGuard 2*, and *OpenAI Moderation API*.

### 3.1 Comparison of Guard Frameworks

System	Dynamic Policy	Logit Threshold	Open Source	Multilingual	Manipulation Defense
<b>OpenGuardrails</b>					
<i>Qwen3Guard</i>	×	×			Partial
<i>LlamaFirewall</i>	×	×	Partial	English-only	
<i>PromptGuard 2</i>	×	×	×	English-only	×
<i>OpenAI Moderation</i>	×	×	×	Limited	×

Table 1: Comparison of key features among major LLM safety and moderation systems. OpenGuardrails uniquely supports both per-request configurable policies and continuous logit-based sensitivity control.

### 3.2 Prior Approaches

**Qwen3Guard** represents one of the earliest multilingual safety guard models, capable of classifying content across 28 unsafe categories. However, its architecture operates in two discrete configurations: *Strict Mode*, which classifies all controversial content as unsafe, and *Loose Mode*, which treats such content as safe. This rigid dichotomy prevents per-category policy customization or continuous sensitivity calibration. Furthermore, Qwen3Guard outputs structured textual results (e.g., “Safety: Safe; Categories: None”), which are semantically interpretable but probabilistically non-differentiable, preventing threshold tuning via token-level logits.

**LlamaFirewall** proposed a two-stage hybrid system combining a large LLM for semantic reasoning with a BERT-based detector (*PromptGuard 2*) for classification. While effective against jailbreaks and prompt injections, the approach suffers from limited multilingual support and high pipeline latency due to multi-model inference. The reliance on discrete classifier thresholds also restricts continuous adaptation to varying sensitivity requirements.

**PromptGuard 2** (integrated within LlamaFirewall) is a fine-tuned RoBERTa-like model designed for prompt and output classification. Although computationally efficient, its embedding-based scoring lacks interpretability and adaptability. It cannot dynamically modify safety criteria per API call, which limits its flexibility for enterprise-scale deployments.

**OpenAI Moderation API** is the most widely adopted commercial moderation tool, offering static models fine-tuned on human-labeled safety data. While highly reliable in English, its performance in low-resource languages remains limited, and it does not expose configuration options for sensitivity or unsafe-category subsets. Users must rely on post-processing logic to approximate threshold tuning.

### 3.3 Distinctive Position of OpenGuardrails

In contrast to these prior systems, **OpenGuardrails** achieves the following distinguishing advantages:

1. **Per-request policy control:** Policy configurations (unsafe categories, sensitivity thresholds) are accepted as runtime API parameters, enabling heterogeneous moderation logic across requests without model reloading.
2. **Continuous probabilistic control:** By operating directly in logit space, the model enables real-valued threshold tuning  $\tau \in [0, 1]$ , allowing smooth precision–recall trade-offs.

3. **Unified LLM architecture:** Both content-safety and model-manipulation detection are handled by the same LLM, unlike multi-model pipelines [6].
4. **Enterprise readiness:** The system integrates with quantized inference (GPTQ 3.3B) and RESTful APIs for production deployment, balancing speed and interpretability.

Overall, OpenGuardrails extends the frontier of guard system design from static, rule-based moderation toward adaptive, mathematically interpretable safety governance—bridging the gap between theoretical safety alignment and practical enterprise deployment.

## 4 Configurable Policy Adaptation

A central innovation of **OpenGuardrails** lies in its *Configurable Policy Adaptation* mechanism, which provides unprecedented flexibility and fine-grained control over safety governance at inference time. Traditional guard models—including *Qwen3Guard*—treat policy as a static configuration, fixed at either training or evaluation stage. In contrast, OpenGuardrails exposes the **safety policy as a runtime parameter** of the detection API, allowing each detection call to operate with a distinct configuration. Users can dynamically specify:

- the exact set of *unsafe categories* to be detected, and
- the *sensitivity threshold* used to interpret model confidence scores.

This design enables per-request customization, which is critical for enterprise-level deployments that must comply with heterogeneous regulatory and cultural standards. For instance, a financial institution may emphasize detection of data leakage and fraudulent advice, while a creative-writing platform may disable political content filtering to preserve artistic freedom. Each request to the OpenGuardrails API can therefore enforce a completely different safety policy, without retraining or restarting the model.

By comparison, **Qwen3Guard** implements only a coarse-grained dual-mode configuration:

- **Strict Mode:** treats all “controversial” samples as unsafe;
- **Loose Mode:** treats all “controversial” samples as safe.

While useful for broad testing, this binary toggle cannot support per-category or per-domain policy differences. In particular, Qwen3Guard’s “controversial” tag represents a lumped uncertainty region that must be either globally accepted or globally rejected. This all-or-nothing mechanism increases manual review costs and prevents automated adaptation in production.

OpenGuardrails solves this limitation through two innovations:

1. **Dynamic Unsafe-Category Selection.** Each inference request can include a JSON or YAML configuration specifying which unsafe categories (e.g., sexual, political, violent, data leakage) are active. The detection pipeline loads the configuration dynamically, adjusting prompt templates and logits interpretation accordingly.
2. **Continuous Sensitivity Thresholding.** Instead of the discrete “strict/loose” switch, OpenGuardrails supports a continuous sensitivity parameter  $\tau \in [0, 1]$ , enabling administrators to calibrate detection aggressiveness in real time.

This design yields both operational and theoretical benefits. Operationally, it allows enterprises to define risk tolerance levels at runtime. Theoretically, it formalizes moderation as a probabilistic hypothesis test, as detailed below.

## 4.1 Mathematical Foundations of Sensitivity Thresholding

The sensitivity threshold mechanism in **OpenGuardrails** is grounded in probabilistic decision theory. The model’s output is interpreted as a binary hypothesis test:

$$H_0 : \text{content is safe}, \quad H_1 : \text{content is unsafe}.$$

Let  $x$  denote the input prompt or response to be evaluated. The guard model defines a conditional probability distribution over output tokens:

$$P_\theta(y_t \mid x, y_{<t}),$$

where  $y_t$  denotes the  $t$ -th generated token. The model’s immediate safety judgment is contained in the first token  $y_1$ , which belongs to the candidate set  $\mathcal{V} = \{\text{safe}, \text{unsafe}\}$ .

The unsafe probability is computed from the softmax-normalized logits of the first token:

$$p_{\text{unsafe}} = \frac{\exp(z_{\text{unsafe}})}{\exp(z_{\text{safe}}) + \exp(z_{\text{unsafe}})},$$

where  $z_i$  denotes the pre-softmax logit value of token  $i$ .

The decision function is defined as:

$$D(x) = \begin{cases} \text{unsafe}, & \text{if } p_{\text{unsafe}} \geq \tau, \\ \text{safe}, & \text{otherwise.} \end{cases}$$

Here,  $\tau \in [0, 1]$  is the configurable **sensitivity threshold**. Lower values of  $\tau$  increase recall by detecting borderline unsafe cases, while higher values increase precision by filtering out low-confidence detections. In practice, OpenGuardrails supports both numeric and semantic mappings, e.g.:

$$\tau_{\text{low}} = 0.3, \quad \tau_{\text{medium}} = 0.5, \quad \tau_{\text{high}} = 0.7.$$

This formulation allows continuous control over the model’s moderation strictness, providing a smooth trade-off between false positives and false negatives.

By contrast, **Qwen3Guard** lacks this probabilistic interface. Its generation output is structured text such as:

```
Safety: Safe
Categories: None
Refusal: Yes
```

Because the model directly emits textual tags, rather than exposing token-level probabilities, the first-token logits are not semantically meaningful for thresholding. The classification decision is encoded in a textual pattern learned during supervised fine-tuning, not in a normalized probabilistic space. As a result, Qwen3Guard cannot perform continuous calibration of sensitivity or per-category adaptation.

Formally, Qwen3Guard’s inference function can be expressed as:

$$D_{\text{Qwen}}(x) = f_\theta(x; m),$$

where  $m \in \{\text{Strict}, \text{Loose}\}$  controls only the high-level controversial toggle. The model outputs deterministic textual labels, and  $\partial D / \partial z_i$  is undefined at inference time, making gradient-based or probabilistic adjustment impossible.

In contrast, OpenGuardrails defines its decision directly in logit space:

$$\frac{\partial D}{\partial z_i} \neq 0,$$

enabling differentiable or probabilistic calibration over token logits. This makes it possible to dynamically adjust safety sensitivity, implement soft decision boundaries, and support domain-dependent detection policies—all within a unified mathematical framework.

## 4.2 Discussion

The combination of per-request policy configuration and probabilistic sensitivity thresholding represents a substantial advancement over existing guard architectures. By elevating policy from a static rule set to a dynamic inference parameter, OpenGuardrails enables fine-grained, mathematically grounded safety governance. This approach transforms safety detection from a binary tagging problem into a tunable probabilistic system, reducing manual review overhead and aligning model behavior with real-world regulatory diversity.

In essence, while Qwen3Guard’s “Strict/Loose” dichotomy provides a static binary choice, OpenGuardrails offers a continuous, configurable, and interpretable control surface over safety sensitivity. This capability is critical for enterprise deployments that demand nuanced, data-driven safety configurations under variable legal, cultural, or domain-specific constraints.

## 5 Evaluation

We evaluate **OpenGuardrails** following the methodology of **Qwen3Guard** to ensure fair comparability.

### English Prompt Results

Table 2: English Prompt Results. Higher is better (F1 score).

Model	Toxic	OpenAIMod	Aegis	Aegis2.0	SimpST	WildG	Avg.
LlamaGuard3-8B	53.8	79.5	71.5	76.4	99.5	76.4	76.2
LlamaGuard4-12B	51.3	73.5	67.8	70.6	98.0	73.0	72.4
WildGuard-7B	70.8	72.1	89.4	80.7	99.5	88.9	83.6
ShieldGemma-9B	69.4	82.1	70.3	72.5	83.7	54.2	72.0
ShieldGemma-27B	72.9	80.5	69.0	71.6	84.4	54.3	72.1
NemoGuard-8B	75.6	81.0	81.4	86.8	98.5	81.6	84.2
PolyGuard-Qwen-7B	71.5	74.1	90.3	86.3	<b>100.0</b>	88.1	85.1
Qwen3Guard-0.6B-Gen (strict)	65.1	66.5	90.8	85.0	99.0	87.7	82.4
Qwen3Guard-0.6B-Gen (loose)	77.7	77.6	76.9	83.3	95.8	85.1	82.7
Qwen3Guard-4B-Gen (strict)	69.5	68.3	90.8	85.8	99.5	85.6	83.3
Qwen3Guard-4B-Gen (loose)	82.8	80.7	76.3	82.1	97.4	85.1	84.1
Qwen3Guard-8B-Gen (strict)	68.9	68.8	<b>91.4</b>	86.1	99.5	<b>88.9</b>	83.9
Qwen3Guard-8B-Gen (loose)	<b>82.8</b>	81.3	76.0	82.5	97.4	85.6	84.3
<b>OpenGuardrails-Text-2510</b>	79.1	<b>86.6</b>	82.0	<b>90.1</b>	98.0	86.8	<b>87.1</b>

## English Response Results

Table 3: English Response Results. Higher is better (F1 score).

Model	SafeRLHF	Beavertails	Aegis2.0	WildG	Avg.
LlamaGuard3-8B	45.2	67.9	66.1	69.5	62.2
LlamaGuard4-12B	42.5	68.6	63.7	66.4	60.3
WildGuard-7B	64.2	84.4	83.2	75.4	76.8
ShieldGemma-9B	44.2	62.4	70.8	49.9	56.8
ShieldGemma-27B	52.6	67.6	74.9	52.4	61.9
NemoGuard-8B	57.6	78.5	<b>87.6</b>	77.5	75.3
PolyGuard-Qwen-7B	63.3	79.5	81.9	77.9	75.7
Qwen3Guard-0.6B-Gen (strict)	66.6	86.1	84.2	76.3	78.3
Qwen3Guard-0.6B-Gen (loose)	64.2	85.4	84.1	77.3	77.8
Qwen3Guard-4B-Gen (strict)	69.8	86.6	86.1	79.5	80.5
Qwen3Guard-4B-Gen (loose)	64.5	85.2	86.5	77.3	78.4
Qwen3Guard-8B-Gen (strict)	70.5	<b>86.6</b>	86.1	78.9	80.5
Qwen3Guard-8B-Gen (loose)	64.2	85.5	86.4	77.3	78.4
<b>OpenGuardrails-Text-2510</b>	<b>95.4</b>	85.3	87.3	<b>85.8</b>	<b>88.5</b>

## Chinese Prompt Results

Table 4: Chinese Prompt Results. Higher is better (F1 score).

Model	ToxiC	WildG	XSTest	Avg.
LlamaGuard3-8B	46.6	70.3	87.9	68.3
LlamaGuard4-12B	47.8	65.6	82.1	65.2
WildGuard-7B	65.6	82.0	83.2	76.9
ShieldGemma-9B	62.8	49.2	78.9	63.6
ShieldGemma-27B	67.2	50.6	80.8	66.2
NemoGuard-8B	51.0	60.7	83.5	65.1
PolyGuard-Qwen-7B	69.7	87.2	54.2	70.4
Qwen3Guard-0.6B-Gen (strict)	64.8	84.8	88.3	79.3
Qwen3Guard-0.6B-Gen (loose)	73.4	83.1	88.5	81.7
Qwen3Guard-4B-Gen (strict)	66.7	87.0	89.4	81.0
Qwen3Guard-4B-Gen (loose)	78.8	84.7	<b>94.1</b>	85.9
Qwen3Guard-8B-Gen (strict)	68.0	<b>88.0</b>	88.2	81.4
Qwen3Guard-8B-Gen (loose)	78.7	84.8	93.3	85.6
<b>OpenGuardrails-Text-2510</b>	<b>84.8</b>	84.7	92.6	<b>87.4</b>



## Chinese Response Results

Table 5: Chinese Response Results. Higher is better (F1 score).

Model	Beavertail	WildG	Avg.
LlamaGuard3-8B	66.1	66.8	66.5
LlamaGuard4-12B	66.8	54.1	60.5
WildGuard-7B	75.4	69.8	72.6
ShieldGemma-9B	59.5	42.8	51.2
ShieldGemma-27B	65.6	47.1	56.4
NemoGuard-8B	72.9	69.4	71.2
PolyGuard-Qwen-7B	79.1	70.2	74.7
Qwen3Guard-0.6B-Gen (strict)	86.2	75.4	80.8
Qwen3Guard-0.6B-Gen (loose)	85.0	73.8	79.4
Qwen3Guard-4B-Gen (strict)	86.7	76.6	81.7
Qwen3Guard-4B-Gen (loose)	84.8	78.2	81.5
Qwen3Guard-8B-Gen (strict)	<b>87.1</b>	77.7	82.4
Qwen3Guard-8B-Gen (loose)	85.1	77.3	81.2
<b>OpenGuardrails</b>	85.6	<b>84.7</b>	<b>85.2</b>

## Multilingual Prompt Results

Table 6: Multilingual Prompt Results on RTP-LX. Higher is better (F1 score).

Model	En	Zh	Ar	Es	Fr	Id	It	Ja	Ko	Ru	Others	Avg.
LlamaGuard3-8B	50.0	47.4	46.6	48.3	49.4	50.7	46.2	49.2	46.6	48.9	46.0	46.6
LlamaGuard4-12B	37.6	35.1	44.4	29.5	36.2	40.7	32.7	50.2	42.1	37.7	42.7	41.7
WildGuard-7B	93.9	80.6	17.3	80.3	74.8	41.5	74.6	53.1	52.9	63.9	37.5	43.9
ShieldGemma-9B	75.8	72.9	50.9	70.6	68.9	61.8	68.3	67.2	65.7	65.0	51.6	55.4
ShieldGemma-27B	76.1	73.4	59.8	71.8	73.6	66.7	73.1	75.0	67.6	74.1	58.2	61.4
NemoGuard-8B	<b>95.4</b>	77.4	21.1	78.1	72.5	34.9	73.4	53.4	67.2	64.0	35.7	42.7
PolyGuard-Qwen-7B	91.2	89.1	84.9	89.0	89.4	74.6	89.3	90.2	86.9	91.3	78.6	80.9
Qwen3Guard-0.6B-Gen (strict)	90.2	85.2	75.7	85.3	87.3	68.2	82.5	87.1	77.1	85.7	72.2	74.8
Qwen3Guard-0.6B-Gen (strict)	90.2	85.2	75.7	85.3	87.3	68.2	85.5	87.9	83.5	89.7	72.9	76.9
Qwen3Guard-0.6B-Gen (loose)	91.4	86.3	76.8	86.5	88.6	69.5	86.9	88.8	84.9	90.2	73.7	77.8
Qwen3Guard-4B-Gen (strict)	92.1	90.6	88.4	88.9	90.8	75.3	88.0	91.3	86.3	91.9	83.9	85.0
Qwen3Guard-4B-Gen (loose)	91.8	90.7	88.3	89.0	90.9	75.5	88.2	91.2	86.2	91.7	83.8	84.9
Qwen3Guard-8B-Gen (strict)	91.5	90.5	88.2	88.7	90.7	75.2	87.8	91.1	86.1	91.6	83.6	84.8
Qwen3Guard-8B-Gen (loose)	91.6	90.6	88.1	88.8	90.8	75.4	87.9	91.2	86.2	91.8	83.7	84.9
<b>OpenGuardrails-Text-2510</b>	90.5	<b>96.2</b>	<b>99.0</b>	<b>98.8</b>	<b>99.0</b>	<b>96.1</b>	<b>99.3</b>	<b>95.4</b>	<b>98.0</b>	<b>99.3</b>	<b>98.6</b>	<b>97.3</b>

## Multilingual Response Results

Table 7: Multilingual Response Results on PolyGuard-Response. Higher is better (F1 score).

Model	En	Zh	Ar	Es	Fr	It	Ja	Ko	Ru	Others	Avg.
LlamaGuard3-8B	69.7	62.8	62.6	67.2	67.1	66.4	65.8	64.0	69.2	65.4	65.8
LlamaGuard4-12B	66.4	56.0	46.8	55.3	55.4	53.3	49.6	51.9	55.5	52.2	53.4
WildGuard-7B	74.5	70.8	44.4	71.7	71.8	71.0	68.0	65.2	71.5	58.8	63.5
ShieldGemma-9B	51.3	46.9	43.6	46.8	49.3	45.9	45.2	44.5	48.8	46.6	46.8
ShieldGemma-27B	53.9	49.9	50.1	48.3	49.9	49.5	51.5	48.2	52.6	50.3	50.4
NemoGuard-8B	76.9	69.0	63.6	72.0	70.2	71.3	65.7	65.9	70.8	69.6	69.5
PolyGuard-Qwen-7B	77.7	70.4	77.2	71.8	72.8	73.1	72.6	73.6	70.4	74.9	74.0
Qwen3Guard-0.6B-Gen (strict)	75.7	74.0	75.8	76.0	74.2	73.9	73.6	75.2	75.9	72.8	74.2
Qwen3Guard-0.6B-Gen (loose)	75.2	75.1	75.4	74.7	74.3	73.7	72.9	74.2	74.9	73.3	74.2
Qwen3Guard-4B-Gen (strict)	79.3	77.6	78.6	79.0	78.7	77.4	76.5	77.6	79.6	77.7	78.1
Qwen3Guard-4B-Gen (loose)	77.7	78.5	78.9	79.1	77.1	77.6	76.4	76.9	79.6	77.7	78.1
Qwen3Guard-8B-Gen (strict)	78.4	76.6	77.2	77.6	78.6	76.9	76.9	77.8	78.2	77.0	77.6
Qwen3Guard-8B-Gen (loose)	78.9	77.1	77.5	78.1	78.8	78.6	76.6	78.6	78.8	77.3	77.6
<b>OpenGuardrails</b>	<b>93.5</b>	<b>96.4</b>	<b>97.8</b>	<b>98.1</b>	<b>98.7</b>	<b>98.1</b>	<b>97.3</b>	<b>96.4</b>	<b>98.5</b>	<b>97.0</b>	<b>97.2</b>

## 6 Discussion

OpenGuardrails demonstrates that it is both technically feasible and operationally practical to deploy a unified LLM-based guard system for real-world enterprise scenarios. The configurable unsafe categories and sensitivity thresholds directly address the *policy inconsistency* problem highlighted in the Qwen3Guard analysis. Instead of relying on a manually labeled “Controversial” class that demands human adjudication, OpenGuardrails enables automated, context-aware safety control through probabilistic sensitivity adjustment. This approach significantly reduces operational overhead while maintaining consistency across diverse regulatory or cultural safety policies.

Moreover, the unified LLM architecture allows OpenGuardrails to capture nuanced semantic patterns of complex prompt-injection or jailbreak attempts—capabilities that smaller classifier-based systems (e.g., BERT-style models in LlamaFirewall) often miss. The deployment-friendly design, combining model quantization and modular API exposure, ensures scalability and low latency in high-concurrency enterprise environments.

Finally, OpenGuardrails sets a new standard for openness in safety infrastructure. By open-sourcing not only the guard model but also a complete production-ready platform and multilingual datasets, it bridges the gap between academic benchmarks and real-world adoption. The project’s demonstrated multilingual coverage and state-of-the-art benchmark results validate its generality and robustness, positioning OpenGuardrails as a foundational framework for trustworthy and secure LLM deployment at scale.

## 7 Limitation

Despite its strong performance, OpenGuardrails still faces several limitations:

1. **Adversarial Robustness.** Although the guardrail model achieves SOTA results in model-manipulation attack detection, it may still be vulnerable to targeted adversarial attacks. Further engineering-based hardening can strengthen the defense.
2. **Fairness and Bias.** Like other moderation systems, fairness and bias in moderation decisions remain open challenges that require continuous evaluation and calibration.

3. **Cross-cultural Adaptation.** Content safety priorities vary across countries and regions beyond universal human values. Region-specific fine-tuning is necessary to adapt to different legal and cultural requirements. Our company provides paid training services to address such localized customization.

## 8 Conclusion

We present OpenGuardrails, an open-source, context-aware guardrails platform with comprehensive coverage across safety, manipulation detection, and data-leak prevention. It outperforms prior systems such as Qwen3Guard, LlamaGuard, and WildGuard in multilingual safety and real-world latency performance.

## References

- [1] Rebedea, T. et al. (2023). NeMo Guardrails: Building Trustworthy LLM Applications. NVIDIA Research.
- [2] AI, G. (2025). RAIL Specification for AI Safety. Guardrails AI.
- [3] Invariant Labs. (2025). Invariant Labs: Open Safety Frameworks for LLMs.
- [4] IBM Research. (2025). Granite Guardian: Safety Inspection for Enterprise AI.
- [5] Qwen Team. *Qwen3Guard Technical Report*. arXiv preprint arXiv:2510.14276, 2025. <http://arxiv.org/abs/2510.14276>.
- [6] Zhang, W., Chen, R., Li, Z., et al. *LlamaFirewall: A Unified Safety and Jailbreak Defense Framework for LLMs*. arXiv preprint arXiv:2403.11211, 2024.
- [7] Hartvigsen, T., Gabriel, S., Holtzman, A., et al. *ToxicChat: A Benchmark for Evaluating and Mitigating LLM Toxicity*. Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP), 2023.
- [8] OpenAI. *OpenAI Moderation API*. 2024. <https://platform.openai.com/docs/guides/moderation>. Accessed: 2025-10-20.
- [9] Li, M., Chen, Q., Zhao, R. *PromptGuard 2: A Lightweight Classifier for Prompt and Output Moderation*. 2024. Included as part of the LlamaFirewall framework.
- [10] Inan, H., Liu, X., Llama Team, et al. *Llama Guard: LLM-based Input-Output Safeguard for Human-AI Conversations*. arXiv preprint arXiv:2312.06674, 2023. <https://arxiv.org/abs/2312.06674>.